

# Enterprise Knowledge Graph Engine

A high-performance, production-ready knowledge graph system designed for complex relationship modeling, semantic reasoning, and large-scale graph analytics. Built for enterprise AI applications requiring scalable, reliable knowledge representation.

## Architecture Overview

The Knowledge Graph Engine provides a comprehensive platform for storing, querying, and analyzing complex relational data through graph structures. It combines the flexibility of NoSQL with the rigor of enterprise database systems.

## Core Components

- **Graph Data Model:** Immutable nodes and edges with rich metadata and versioning
- **Schema Management:** Type-safe validation with constraint enforcement
- **High-Performance Indexing:** Multiple indexing strategies for optimal query performance
- **Advanced Query Engine:** Graph algorithms and pattern matching capabilities
- **Analytics Framework:** Centrality measures, community detection, and path analysis
- **Enterprise Features:** ACID compliance, audit trails, and performance monitoring

## Key Features

### Performance & Scalability

- **In-memory processing** with optional persistence layers
- **Parallel query execution** using thread pools
- **Advanced indexing** with hash, spatial, and semantic indices
- **Optimized graph algorithms** for large-scale analytics
- **Memory-efficient data structures** with automatic cleanup

### Enterprise Reliability

- **Schema validation** with type checking and constraints
- **Versioning system** for audit trails and rollback capabilities
- **ACID compliance** for data consistency guarantees
- **Performance monitoring** with detailed metrics collection
- **Comprehensive error handling** with graceful degradation

### Advanced Analytics

- **Graph algorithms:** Shortest path, centrality measures, community detection
- **Pattern matching:** Complex relationship queries and traversals
- **Semantic reasoning:** Support for inference and rule-based reasoning
- **Machine learning integration:** Embedding support and feature extraction
- **Real-time analytics:** Streaming updates and incremental computation

## Installation and Setup

### Requirements

```
bash
```

```
# Core dependencies
```

```
pip install numpy>=1.21.0
```

```
pip install pandas>=1.3.0
```

```
pip install pydantic>=1.8.0
```

```
# Optional ML dependencies
```

```
pip install scikit-learn>=1.0.0
```

```
pip install networkx>=2.6.0
```

### Basic Setup

```
python
```

```
from enterprise_knowledge_graph import (  
    KnowledgeGraph, GraphNode, GraphEdge,  
    NodeType, EdgeType, GraphSchema  
)
```

```
# Initialize knowledge graph
```

```
config = {  
    'max_workers': 8,  
    'enable_performance_monitoring': True,  
    'cache_size': 10000  
}
```

```
kg = KnowledgeGraph(config)
```

## Quick Start Guide

### 1. Creating Nodes and Edges

python

```
import asyncio
import datetime
from enterprise_knowledge_graph import *
```

```
async def basic_example():
```

```
    # Initialize knowledge graph
```

```
    kg = KnowledgeGraph()
```

```
    # Create nodes
```

```
    person_node = GraphNode(
```

```
        id="person_001",
```

```
        type=NodeType.PERSON,
```

```
        label="John Doe",
```

```
        properties={
```

```
            "age": 35,
```

```
            "occupation": "Data Scientist",
```

```
            "location": "San Francisco",
```

```
            "skills": ["Python", "Machine Learning", "Graph Theory"]
```

```
        }
```

```
    )
```

```
    company_node = GraphNode(
```

```
        id="company_001",
```

```
        type=NodeType.ORGANIZATION,
```

```
        label="TechCorp Inc",
```

```
        properties={
```

```
            "industry": "Technology",
```

```
            "founded": 2010,
```

```
            "employees": 5000,
```

```
            "revenue": 1000000000
```

```
        }
```

```
    )
```

```
    project_node = GraphNode(
```

```
        id="project_001",
```

```
        type=NodeType.ENTITY,
```

```
        label="AI Knowledge System",
```

```
        properties={
```

```
            "status": "active",
```

```
            "budget": 2000000,
```

```
            "deadline": "2024-12-31"
```

```
        }
```

```
    )
```

```
    # Add nodes to graph
```

```

# Add nodes to graph
await kg.add_node(person_node)
await kg.add_node(company_node)
await kg.add_node(project_node)

# Create relationships
works_for_edge = GraphEdge(
    id="edge_001",
    source_id="person_001",
    target_id="company_001",
    type=EdgeType.RELATED_TO,
    label="works_for",
    properties={
        "start_date": "2020-01-15",
        "position": "Senior Data Scientist",
        "salary": 150000
    },
    weight=1.0
)

leads_project_edge = GraphEdge(
    id="edge_002",
    source_id="person_001",
    target_id="project_001",
    type=EdgeType.RELATED_TO,
    label="leads",
    properties={
        "role": "Technical Lead",
        "responsibility_level": "high"
    },
    weight=0.9
)

# Add edges to graph
await kg.add_edge(works_for_edge)
await kg.add_edge(leads_project_edge)

print(f"Graph created with {len(kg.nodes)} nodes and {len(kg.edges)} edges")

# Search for nodes
tech_companies = await kg.search_nodes({
    "type": "organization",
    "industry": "Technology"
})

print(f"Found {len(tech_companies)} technology companies")

```

```
await kg.close()
```

```
asyncio.run(basic_example())
```

## 2. Schema Definition and Validation

python

```
async def schema_example():
    kg = KnowledgeGraph()

    # Define schema for person nodes
    kg.schema.define_node_schema(
        NodeType.PERSON,
        required_properties=["name", "age"],
        property_types={
            "name": str,
            "age": int,
            "email": str
        },
        constraints=[
            lambda node: node.properties.get("age", 0) >= 0,
            lambda node: "@" in node.properties.get("email", "@")
        ]
    )

    # Define schema for employment edges
    kg.schema.define_edge_schema(
        EdgeType.RELATED_TO,
        allowed_source_types=[NodeType.PERSON],
        allowed_target_types=[NodeType.ORGANIZATION],
        required_properties=["start_date"],
        property_types={
            "start_date": str,
            "salary": int
        }
    )

    # Create valid node
    valid_person = GraphNode(
        id="person_002",
        type=NodeType.PERSON,
        label="Jane Smith",
        properties={
            "name": "Jane Smith",
            "age": 28,
            "email": "jane@example.com"
        }
    )

    await kg.add_node(valid_person) # Will succeed

    # Try to create invalid node (will raise validation error)
```

```
# Try to create invalid node (will raise validation error)
try:
    invalid_person = GraphNode(
        id="person_003",
        type=NodeType.PERSON,
        label="Invalid Person",
        properties={
            "name": "Bob",
            "age": -5, # Invalid age
            "email": "invalid-email" # Invalid email
        }
    )
    await kg.add_node(invalid_person)
except ValueError as e:
    print(f"Validation failed: {e}")

await kg.close()

asyncio.run(schema_example())
```

### 3. Advanced Querying and Graph Algorithms

python

```
async def advanced_querying_example():
    kg = KnowledgeGraph()

    # Create a more complex graph
    # ... add nodes and edges (company hierarchy, projects, skills, etc.)

    # Populate sample data
    await create_sample_corporate_graph(kg)

    # Find shortest path between two people
    path = await kg.find_shortest_path("person_001", "person_005")
    print(f"Shortest path: {path}")

    # Calculate centrality measures
    centrality = await kg.calculate_centrality("betweenness")
    most_central = max(centrality.items(), key=lambda x: x[1])
    print(f"Most central node: {most_central[0]} (score: {most_central[1]:.3f})")

    # Detect communities
    communities = await kg.detect_communities("louvain")
    community_sizes = {}
    for node_id, community in communities.items():
        community_sizes[community] = community_sizes.get(community, 0) + 1

    print(f"Detected {len(community_sizes)} communities:")
    for community, size in community_sizes.items():
        print(f"  Community {community}: {size} members")

    # Complex queries
    senior_people = await kg.search_nodes({
        "type": "person",
        "seniority_level": "senior"
    })

    # Find all projects led by senior people
    senior_projects = []
    for person in senior_people:
        outgoing_edges = await kg.get_outgoing_edges(person.id)
        for edge in outgoing_edges:
            if edge.label == "leads" and edge.properties.get("type") == "project":
                project = await kg.get_node(edge.target_id)
                if project:
                    senior_projects.append(project)

    print(f"Projects led by senior people: {len(senior_projects)}")
```



```
print(f' Projects led by senior people: {len(senior_projects)}' )
```

```
await kg.close()
```

```
async def create_sample_corporate_graph(kg):
```

```
    """Create a sample corporate graph structure."""
```

```
    # Create employees
```

```
    employees = [
```

```
        GraphNode(
```

```
            id=f"person_{i:03d}",
```

```
            type=NodeType.PERSON,
```

```
            label=f"Employee {i}",
```

```
            properties={
```

```
                "department": ["Engineering", "Marketing", "Sales", "HR"][i % 4],
```

```
                "seniority_level": ["junior", "mid", "senior"][i % 3],
```

```
                "years_experience": i % 10 + 1
```

```
            }
```

```
        )
```

```
        for i in range(1, 21) # 20 employees
```

```
    ]
```

```
    # Create departments
```

```
    departments = [
```

```
        GraphNode(
```

```
            id="dept_eng",
```

```
            type=NodeType.ORGANIZATION,
```

```
            label="Engineering",
```

```
            properties={"budget": 5000000, "headcount": 50}
```

```
        ),
```

```
        GraphNode(
```

```
            id="dept_marketing",
```

```
            type=NodeType.ORGANIZATION,
```

```
            label="Marketing",
```

```
            properties={"budget": 2000000, "headcount": 20}
```

```
        )
```

```
    ]
```

```
    # Create projects
```

```
    projects = [
```

```
        GraphNode(
```

```
            id=f"project_{i:03d}",
```

```
            type=NodeType.ENTITY,
```

```
            label=f"Project {i}",
```

```
            properties={
```

```
                "status": ["active", "completed", "on_hold"][i % 3],
```

```
                "priority": ["high", "medium", "low"][i % 3]
```

```
            }
```

```

    )
    for i in range(1, 11) # 10 projects
]

# Add all nodes
for node in employees + departments + projects:
    await kg.add_node(node)

# Create relationships
edge_id = 1

# Employee -> Department relationships
for emp in employees:
    dept_id = "dept_eng" if emp.properties["department"] == "Engineering" else "dept_mkt"
    edge = GraphEdge(
        id=f"edge_{edge_id:03d}",
        source_id=emp.id,
        target_id=dept_id,
        type=EdgeType.PART_OF,
        label="member_of",
        properties={"join_date": "2020-01-01"}
    )
    await kg.add_edge(edge)
    edge_id += 1

# Employee -> Project relationships
for i, project in enumerate(projects):
    # Assign 2-3 people per project
    assigned_people = employees[i*2:(i*2)+3]
    for j, person in enumerate(assigned_people):
        edge = GraphEdge(
            id=f"edge_{edge_id:03d}",
            source_id=person.id,
            target_id=project.id,
            type=EdgeType.RELATED_TO,
            label="works_on" if j > 0 else "leads",
            properties={"allocation": 0.5 if j > 0 else 1.0}
        )
        await kg.add_edge(edge)
        edge_id += 1

asyncio.run(advanced_querying_example())

```

## 4. Graph Analytics and Machine Learning Integration

python

```
import numpy as np
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA

async def ml_integration_example():
    kg = KnowledgeGraph()

    # Create sample graph with embeddings
    await create_sample_graph_with_embeddings(kg)

    # Extract node embeddings for analysis
    embeddings = []
    node_ids = []

    for node_id, node in kg.nodes.items():
        if node.embedding is not None:
            embeddings.append(node.embedding)
            node_ids.append(node_id)

    if embeddings:
        embeddings_array = np.array(embeddings)

        # Perform clustering on embeddings
        kmeans = KMeans(n_clusters=3, random_state=42)
        clusters = kmeans.fit_predict(embeddings_array)

        # Visualize with PCA
        pca = PCA(n_components=2)
        embeddings_2d = pca.fit_transform(embeddings_array)

        print("Node clustering based on embeddings:")
        for i, (node_id, cluster) in enumerate(zip(node_ids, clusters)):
            node = kg.nodes[node_id]
            print(f"{node.label}: Cluster {cluster}")

        # Find similar nodes using embedding distance
        def find_similar_nodes(target_node_id, top_k=3):
            target_embedding = kg.nodes[target_node_id].embedding
            similarities = []

            for node_id, node in kg.nodes.items():
                if node_id != target_node_id and node.embedding is not None:
                    # Cosine similarity
                    similarity = np.dot(target_embedding, node.embedding) / (
                        np.linalg.norm(target_embedding) * np.linalg.norm(node.embedding))
```

```

        np.linalg.norm(target_embedding) * np.linalg.norm(node_embedding)
    )
    similarities.append((node_id, similarity))

    similarities.sort(key=lambda x: x[1], reverse=True)
    return similarities[:top_k]

# Find nodes similar to the first node
if node_ids:
    similar = find_similar_nodes(node_ids[0])
    print(f"\nNodes similar to {kg.nodes[node_ids[0]].label}:")
    for node_id, similarity in similar:
        print(f"    {kg.nodes[node_id].label}: {similarity:.3f}")

await kg.close()

async def create_sample_graph_with_embeddings(kg):
    """Create sample graph with node embeddings."""
    categories = ["technology", "science", "arts", "business"]

    for i in range(20):
        # Generate random embedding
        embedding = np.random.normal(0, 1, 128)

        # Add category bias to embedding
        category = categories[i % len(categories)]
        if category == "technology":
            embedding[:32] += 2.0
        elif category == "science":
            embedding[32:64] += 2.0
        elif category == "arts":
            embedding[64:96] += 2.0
        else: # business
            embedding[96:] += 2.0

        node = GraphNode(
            id=f"concept_{i:03d}",
            type=NodeType.CONCEPT,
            label=f"{category.title()} Concept {i}",
            properties={
                "category": category,
                "importance": np.random.random()
            },
            embedding=embedding
        )

        await kg.add_node(node)

```

```
asyncio.run(ml_integration_example())
```

## Production Deployment

### High-Performance Configuration

python

```
# production_config.py
from enterprise_knowledge_graph import KnowledgeGraph

def create_production_kg():
    """Create production-optimized knowledge graph."""
    config = {
        # Performance settings
        'max_workers': 16,
        'enable_performance_monitoring': True,
        'cache_size': 100000,

        # Memory management
        'memory_limit_gb': 8,
        'gc_threshold': 10000,

        # Persistence settings
        'auto_save_interval': 300, # 5 minutes
        'backup_retention_days': 30,

        # Indexing strategy
        'indexing_strategy': 'hash',
        'enable_spatial_index': True,
        'enable_text_index': True,

        # Query optimization
        'query_cache_size': 10000,
        'max_query_depth': 20,
        'query_timeout_seconds': 30
    }

    return KnowledgeGraph(config)
```

### Persistence and Backup

python

```
import json
import pickle
from pathlib import Path

class KnowledgeGraphPersistence:
    """Handles persistence and backup for knowledge graphs."""

    def __init__(self, kg: KnowledgeGraph, storage_path: Path):
        self.kg = kg
        self.storage_path = Path(storage_path)
        self.storage_path.mkdir(parents=True, exist_ok=True)

    async def save_to_json(self, filename: str = None):
        """Save graph to JSON format."""
        if filename is None:
            timestamp = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")
            filename = f"knowledge_graph_{timestamp}.json"

        filepath = self.storage_path / filename
        graph_data = await self.kg.export_to_dict()

        with open(filepath, 'w') as f:
            json.dump(graph_data, f, indent=2, default=str)

        return filepath

    async def load_from_json(self, filepath: Path):
        """Load graph from JSON format."""
        with open(filepath, 'r') as f:
            graph_data = json.load(f)

        await self.kg.import_from_dict(graph_data)
        return True

    async def save_to_pickle(self, filename: str = None):
        """Save graph to pickle format (faster but less portable)."""
        if filename is None:
            timestamp = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")
            filename = f"knowledge_graph_{timestamp}.pkl"

        filepath = self.storage_path / filename
        graph_data = await self.kg.export_to_dict()

        with open(filepath, 'wb') as f:
            pickle.dump(graph_data, f)
```

```
pickle.dump(graph_data, f)
```

```
return filepath
```

```
async def create_backup(self):
```

```
    """Create timestamped backup."""
```

```
    timestamp = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")
```

```
    backup_dir = self.storage_path / "backups" / timestamp
```

```
    backup_dir.mkdir(parents=True, exist_ok=True)
```

```
    # Save in multiple formats
```

```
    json_path = await self.save_to_json(str(backup_dir / "graph.json"))
```

```
    pickle_path = await self.save_to_pickle(str(backup_dir / "graph.pkl"))
```

```
    # Save metadata
```

```
    metadata = {
```

```
        "created_at": timestamp,
```

```
        "statistics": self.kg.get_statistics(),
```

```
        "formats": ["json", "pickle"]
```

```
    }
```

```
    with open(backup_dir / "metadata.json", 'w') as f:
```

```
        json.dump(metadata, f, indent=2, default=str)
```

```
    return backup_dir
```

```
# Example usage
```

```
async def persistence_example():
```

```
    kg = KnowledgeGraph()
```

```
    # Create sample data
```

```
    await create_sample_corporate_graph(kg)
```

```
    # Setup persistence
```

```
    persistence = KnowledgeGraphPersistence(kg, Path("./kg_storage"))
```

```
    # Save graph
```

```
    json_path = await persistence.save_to_json()
```

```
    print(f"Graph saved to: {json_path}")
```

```
    # Create backup
```

```
    backup_dir = await persistence.create_backup()
```

```
    print(f"Backup created: {backup_dir}")
```

```
    # Load into new graph
```

```
    new_kg = KnowledgeGraph()
```

```
    new_persistence = KnowledgeGraphPersistence(new_kg, Path("./kg_storage"))
```

```
    await new_persistence.load_from_json(json_path)

    print(f"Loaded graph: {len(new_kg.nodes)} nodes, {len(new_kg.edges)} edges")

    await kg.close()
    await new_kg.close()

asyncio.run(persistence_example())
```

## Distributed Processing



python

```
import concurrent.futures
from typing import List, Callable

class DistributedGraphProcessor:
    """Handles distributed processing of large graphs."""

    def __init__(self, kg: KnowledgeGraph, num_workers: int = 4):
        self.kg = kg
        self.num_workers = num_workers
        self.executor = concurrent.futures.ThreadPoolExecutor(max_workers=num_workers)

    async def parallel_node_processing(
        self,
        processor_func: Callable[[GraphNode], Any],
        batch_size: int = 1000
    ) -> List[Any]:
        """Process nodes in parallel batches."""
        nodes = list(self.kg.nodes.values())
        results = []

        # Process in batches
        for i in range(0, len(nodes), batch_size):
            batch = nodes[i:i + batch_size]

            # Submit batch to thread pool
            futures = [
                self.executor.submit(processor_func, node)
                for node in batch
            ]

            # Collect results
            batch_results = [future.result() for future in futures]
            results.extend(batch_results)

            print(f"Processed batch {i//batch_size + 1}/{(len(nodes)-1)//batch_size + 1}")

        return results

    async def parallel_subgraph_analysis(
        self,
        analyzer_func: Callable[[List[GraphNode], List[GraphEdge]], Any],
        partition_size: int = 500
    ) -> List[Any]:
        """Analyze graph partitions in parallel."""
        # Simple partitioning by node ID hash
```

```

# Simple partitioning by node ID hash
partitions = {}

for node in self.kg.nodes.values():
    partition_id = hash(node.id) % self.num_workers
    if partition_id not in partitions:
        partitions[partition_id] = {'nodes': [], 'edges': []}
    partitions[partition_id]['nodes'].append(node)

# Add edges to appropriate partitions
for edge in self.kg.edges.values():
    source_partition = hash(edge.source_id) % self.num_workers
    partitions[source_partition]['edges'].append(edge)

# Process partitions in parallel
futures = [
    self.executor.submit(
        analyzer_func,
        partition['nodes'],
        partition['edges']
    )
    for partition in partitions.values()
]

results = [future.result() for future in futures]
return results

def close(self):
    """Clean up thread pool."""
    self.executor.shutdown(wait=True)

# Example usage
async def distributed_processing_example():
    kg = KnowledgeGraph()
    await create_sample_corporate_graph(kg)

    processor = DistributedGraphProcessor(kg, num_workers=4)

# Example: Extract features from all nodes in parallel
def extract_features(node: GraphNode) -> Dict[str, Any]:
    return {
        'id': node.id,
        'type': node.type.value,
        'property_count': len(node.properties),
        'has_embedding': node.embedding is not None
    }

```

```

features = await processor.parallel_node_processing(extract_features)
print(f"Extracted features from {len(features)} nodes")

# Example: Analyze subgraphs in parallel
def analyze_subgraph(nodes: List[GraphNode], edges: List[GraphEdge]) -> Dict[str, float]:
    return {
        'node_count': len(nodes),
        'edge_count': len(edges),
        'density': len(edges) / (len(nodes) * (len(nodes) - 1)) if len(nodes) > 1 else 0
    }

subgraph_analyses = await processor.parallel_subgraph_analysis(analyze_subgraph)
print(f"Analyzed {len(subgraph_analyses)} subgraphs")

processor.close()
await kg.close()

asyncio.run(distributed_processing_example())

```

## Performance Monitoring and Optimization

### Metrics Collection

python

```
class GraphMetrics:
    """Advanced metrics collection for knowledge graphs."""

    def __init__(self, kg: KnowledgeGraph):
        self.kg = kg
        self.start_time = datetime.datetime.now()

    def get_comprehensive_metrics(self) -> Dict[str, Any]:
        """Get comprehensive performance metrics."""
        stats = self.kg.get_statistics()

        # Calculate additional metrics
        uptime = (datetime.datetime.now() - self.start_time).total_seconds()

        # Graph structure metrics
        avg_degree = (2 * stats['edges']) / max(1, stats['nodes'])
        density = (2 * stats['edges']) / max(1, stats['nodes'] * (stats['nodes'] - 1))

        # Performance metrics
        queries_per_second = stats['metrics']['queries_executed'] / max(1, uptime)

        return {
            **stats,
            'uptime_seconds': uptime,
            'average_degree': avg_degree,
            'graph_density': density,
            'queries_per_second': queries_per_second,
            'memory_efficiency': stats['memory_usage_bytes'] / max(1, stats['nodes'] +
        }

    async def benchmark_operations(self, iterations: int = 1000) -> Dict[str, float]:
        """Benchmark common graph operations."""
        import time

        # Node addition benchmark
        start_time = time.time()
        for i in range(iterations):
            node = GraphNode(
                id=f"benchmark_node_{i}",
                type=NodeType.ENTITY,
                label=f"Benchmark Node {i}",
                properties={"benchmark": True}
            )
            await self.kg.add_node(node)
```

```

node_addition_time = (time.time() - start_time) / iterations

# Query benchmark
start_time = time.time()
for i in range(iterations):
    await self.kg.search_nodes({"benchmark": True}, limit=10)

query_time = (time.time() - start_time) / iterations

# Cleanup benchmark nodes
benchmark_nodes = await self.kg.search_nodes({"benchmark": True})
for node in benchmark_nodes:
    await self.kg.remove_node(node.id)

return {
    'node_addition_ms': node_addition_time * 1000,
    'query_ms': query_time * 1000,
    'operations_per_second': 1 / max(node_addition_time, 1e-6)
}

# Example usage
async def monitoring_example():
    kg = KnowledgeGraph()
    await create_sample_corporate_graph(kg)

    metrics_collector = GraphMetrics(kg)

    # Get comprehensive metrics
    metrics = metrics_collector.get_comprehensive_metrics()
    print("Graph Metrics:")
    for key, value in metrics.items():
        if isinstance(value, float):
            print(f"  {key}: {value:.3f}")
        else:
            print(f"  {key}: {value}")

    # Run performance benchmark
    benchmark_results = await metrics_collector.benchmark_operations(100)
    print("\nBenchmark Results:")
    for operation, time_ms in benchmark_results.items():
        print(f"  {operation}: {time_ms:.3f}")

    await kg.close()

asyncio.run(monitoring_example())

```

# Best Practices

## 1. Schema Design

- **Define schemas early** for data consistency and validation
- **Use appropriate node types** to enable type-specific queries
- **Implement constraints** to maintain data quality
- **Version your schemas** for backward compatibility

## 2. Performance Optimization

- **Use appropriate indexing** strategies for your query patterns
- **Batch operations** when adding large amounts of data
- **Monitor memory usage** and implement cleanup procedures
- **Optimize query patterns** by analyzing execution times

## 3. Data Modeling

- **Keep properties lightweight** to improve memory efficiency
- **Use embeddings** for semantic similarity and ML integration
- **Model temporal relationships** explicitly when needed
- **Implement proper versioning** for audit trails

## 4. Production Deployment

- **Implement proper backup strategies** with multiple formats
- **Monitor performance metrics** continuously
- **Use distributed processing** for large-scale analytics
- **Implement proper error handling** and graceful degradation

## Integration Patterns

### Database Integration

python

*# Example: Sync with relational database*

```
async def sync_with_database(kg: KnowledgeGraph, db_connection):
    """Sync knowledge graph with relational database."""

    # Load entities from database
    entities = await db_connection.fetch("SELECT * FROM entities")

    for entity in entities:
        node = GraphNode(
            id=f"entity_{entity['id']}",
            type=NodeType.ENTITY,
            label=entity['name'],
            properties={
                'db_id': entity['id'],
                'created_at': entity['created_at'],
                'updated_at': entity['updated_at']
            }
        )
        await kg.add_node(node)

    # Load relationships
    relationships = await db_connection.fetch("SELECT * FROM relationships")

    for rel in relationships:
        edge = GraphEdge(
            id=f"rel_{rel['id']}",
            source_id=f"entity_{rel['source_id']}",
            target_id=f"entity_{rel['target_id']}",
            type=EdgeType.RELATED_TO,
            label=rel['relationship_type'],
            properties={'db_id': rel['id']}
        )
        await kg.add_edge(edge)
```

## API Integration

python

*# Example: REST API for knowledge graph*

```
from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
```

```
app = FastAPI()
kg = KnowledgeGraph()
```

```
class NodeCreate(BaseModel):
    id: str
    type: str
    label: str
    properties: Dict[str, Any] = {}
```

```
class EdgeCreate(BaseModel):
    id: str
    source_id: str
    target_id: str
    type: str
    label: str
    properties: Dict[str, Any] = {}
```

```
@app.post("/nodes")
async def create_node(node_data: NodeCreate):
    try:
        node = GraphNode(
            id=node_data.id,
            type=NodeType(node_data.type),
            label=node_data.label,
            properties=node_data.properties
        )
        await kg.add_node(node)
        return {"status": "success", "id": node.id}
    except Exception as e:
        raise HTTPException(status_code=400, detail=str(e))
```

```
@app.get("/nodes/{node_id}")
async def get_node(node_id: str):
    node = await kg.get_node(node_id)
    if not node:
        raise HTTPException(status_code=404, detail="Node not found")
    return node.to_dict()
```

```
@app.get("/search/nodes")
async def search_nodes(query: str, limit: int = 100):
    # Parse query string into search parameters
```



```
# Parse query string into search parameters  
search_params = {"label": query} # Simplified  
results = await kg.search_nodes(search_params, limit)  
return [node.to_dict() for node in results]
```

## License

This Knowledge Graph Engine is released under the MIT License, making it suitable for both commercial and open-source AI projects.

## Contributing

Contributions are welcome! Areas of particular interest:

- Additional graph algorithms
- New indexing strategies
- Performance optimizations
- Integration patterns
- Documentation improvements

For questions or support, please refer to the project documentation or open an issue in the repository.